

How to Evolve a Knowledge Graph?



Outline

- Overview
- Examples requiring Change
- Change Management Techniques
 - Schema evolution
 - View maintenance
 - Truth Maintenance
- Summary



Overview



Only constant in life is change

Change in the real-world
Change in the business requirements



Changes can require

Revising the schema
Revising the ground facts



Approaches to handle change
must address

Technical challenges
Social challenges



Examples Requiring Change

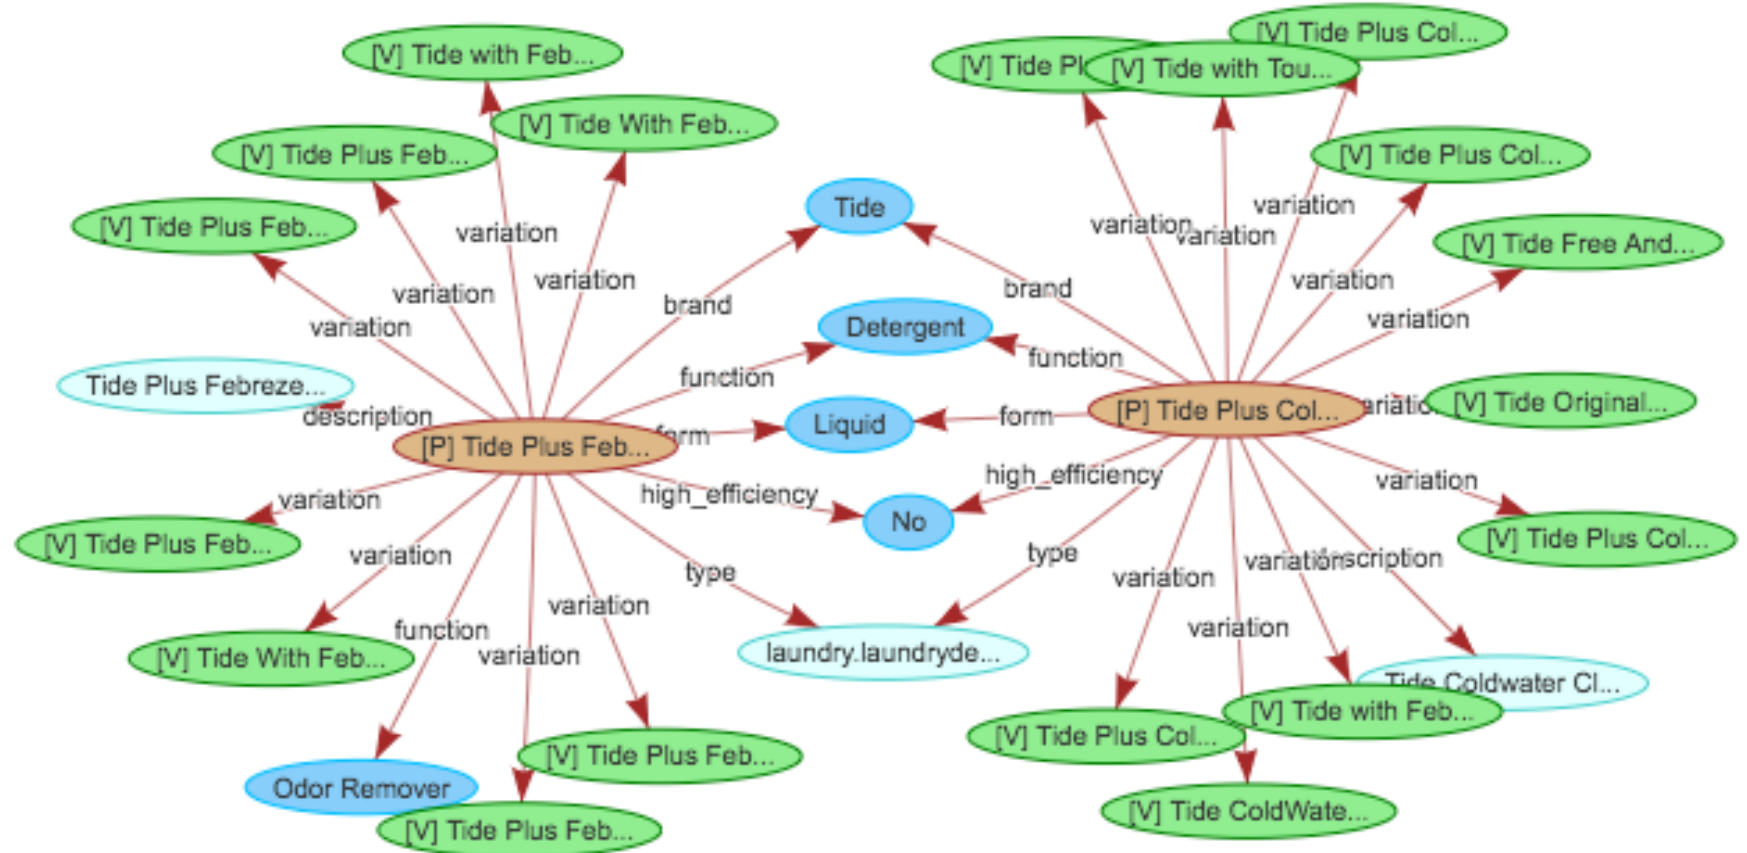
- Changing world
- Changing requirements
- Changing sources
- Changes affecting previous inferences
- Changes requiring redesign



Examples Requiring Change

- Changing world: Amazon Product Knowledge Graph

New products
New product categories
New features
Discontinued products



Examples Requiring Change

- Changing requirements: Google Knowledge Graph

An artist must be a person



An artist must be a person OR a vocaloid



Examples Requiring Change

- Changing sources (Google Knowledge Graph)
 - The artists of music albums are obtained from different sources
 - These sources keep changing their data feed
 - The sources used also change over a period of time




Examples Requiring Change

- Changes affecting previous inferences
 - Consider the constraint that a movie theater only shows movies
 - Using this constraint a KG might have previously inferred that certain events are movies
 - More recently the movie theaters are being used for operas, and social events
 - If we had previously derived such events to be movies, we must update them



Examples Requiring Change


- Changes requiring redesign
 - Initial design: Every company has a CEO. A CEO is represented by name.




Examples Requiring Change

- Changes requiring redesign

- Initial design: Every company has a CEO. A CEO is represented by name.



- Revised design: Every company has a CEO. A CEO is represented by an object that can also record the duration for which the person was a CEO



Outline

- Overview
- Examples requiring Change
- Change Management Techniques
 - Schema evolution
 - View maintenance
 - Truth Maintenance
- Summary



Schema Evolution

- For a relational database
 - Adding/removing a column, renaming an attribute
 - Known as database reorganization



Schema Evolution

- For a relational database
 - Adding/removing a column, renaming an attribute
 - Known as database reorganization
- For a knowledge graph
 - Adding/removing a class
 - Adding/removing a superclass
 - Adding/removing a property
 - Adding/removing a constraint

Approach is to maintain invariants, and make system-specific decisions



Schema Evolution

- Remove/rename a property
 - The change must be propagated through the graph
 - Generate a summary for review by the user




Schema Evolution

- Add a class
 - By default, its parent is the system defined root class




Schema Evolution

- Remove a class
 - What to do about its subclasses and instances
 - If its subclass has another parent, do nothing
 - Otherwise, make it a subclass of immediate parent
 - A more extreme step is to delete the subclasses and instances




Schema Evolution

- Remove a class
 - What to do about its subclasses and instances
 - If its subclass has another parent, do nothing
 - Otherwise, make it a subclass of immediate parent
 - A more extreme step is to delete the subclasses and instances




Schema Evolution

- Remove a class
 - What to do about its subclasses and instances
 - If its subclass has another parent, do nothing
 - Otherwise, make it a subclass of immediate parent
 - A more extreme step is to delete the subclasses and instances




Schema Evolution

- Remove a class
 - What to do about its subclasses and instances
 - If its subclass has another parent, do nothing
 - Otherwise, make it a subclass of immediate parent
 - A more extreme step is to delete the subclasses and instances



Schema Evolution

- Add a super class
 - How to handle redundant links?
 - How to handle cycles?



View Maintenance

- A mechanism from databases to name a query
 - Query is defined with respect to one or more tables (known as **base tables**)
 - If we store the results of the query, the stored data is called **materialized** view
- If the base data changes, the materialized view must be updated
 - Incremental view maintenance



View Maintenance

- A mechanism from databases to name a query
 - Query is defined with respect to one or more tables (known as **base tables**)
 - If we store the results of the query, the stored data is called **materialized** view
- If the base data changes, the materialized view must be updated
 - Incremental view maintenance

Use of view maintenance is not prevalent in current Knowledge Graph engines



Truth Maintenance

- A mechanism from rule-based systems
 - Tracks how each conclusion was derived
- A popular implementation: Justification based system
 - Each derived conclusion records the fact or rule that was used in derivation
 - Any time that fact or rule updates, the conclusion must be revised

Use of truth maintenance is not prevalent in current Knowledge Graph engines



Summary

- Knowledge Graphs have a life-cycle
 - Must evolve over a period of time
 - Must address both social and technical concerns
- Techniques and algorithms
 - Schema evolution
 - View maintenance
 - Truth maintenance

